

# LING 572 Hw2

## Due: 11pm on Jan 18, 2022

For all the assignments in ling572:

- If your code is slow, do not run your code on patas directly. Use “condor submit” instead (see Q5 below).
- If patas is down, you can try dryas (e.g., ssh userid@dryas.ling.washington.edu).
- You can assume that all the class labels in the test data have appeared in the training data; that is, the test data does not contain any new class labels.

The example files for hw2 are under `/dropbox/21-22/572/hw2/examples/`.

**Q1 (4 points):** Run the Mallet DT learner (i.e., the trainer’s name is `DecisionTree`) with **train.vectors.txt** as the training data and **test.vectors.txt** as the test data. In your note file, write down the following:

- (a) The command lines you use for preparing data, training, testing, and getting the training and test accuracy. You can use `vectors2classify` commands to do training, testing and evaluation in one step.
- (b) What are the training accuracy and the test accuracy?

**Q2 (6 points):** Run the Mallet DT trainer with different depths; that is, when running `vectors2classify`, replace `-trainer DecisionTree` with

```
--trainer "new DecisionTreeTrainer(nn)"
```

where `nn` is the depth of the decision tree. Note that you have to use `vectors2classify`, instead of “mallet train-classifier” and “mallet classify-svmlight” because “mallet train-classifier” does not process “new `DecisionTreeTrainer(nn)`” properly.

- (a) Fill out Table 1
- (b) What conclusion can you draw from Table 1?

**Q3 (55 points):** Write a program, **build\_dt.sh**, that builds a DT tree from the training data, classifies the training and test data, and calculates the accuracy.

- This DT learner should treat all features as binary; that is, the feature is considered present if its value is nonzero, and absent if its value is zero.

Table 1: Run Mallet’s DT learner with different depths

| Depth | Training accuracy | Test accuracy |
|-------|-------------------|---------------|
| 1     |                   |               |
| 2     |                   |               |
| 4     |                   |               |
| 10    |                   |               |
| 20    |                   |               |
| 50    |                   |               |

- Use information gain to select features when building DT.
- The format of the command line would be: `build_dt.sh training_data test_data max_depth min_gain model_file sys_output > acc_file`
- `training_data` and `test_data` are the vector files in the text format (cf. **train.vectors.txt**).
- `max_depth` is the maximum depth of the DT,<sup>1</sup> and `min_gain` is the minimal gain. Those parameters are used to determine when to stop building DT; that is, split the current training data set at the node `x` if and only if **(the depth of `x` < `max_depth`) AND (the infoGain of the split  $\geq$  `min_gain`)**.
- `model_file` is the DT tree (cf. **model.ex**) produced by the DT trainer. Each line corresponds to a leaf node in the DT and it has the format: `path training_instance_num c1 p1 c2 p2 ...`. Where `path` is the path from the root to the leaf node, `training_instance_num` is the number of the training examples that “reach” the leaf node, `ci` is the class label, and `pi` is the probability of `ci` (i.e., the percentage of the training examples at the leaf node with the label `ci`).
- `sys_output` is the classification result on the training and test data (cf. **sys.ex**). Each line has the following format:  
`instanceName c1 p1 c2 p2 ...`, where `instanceName` is just something like “array:0”, “array:1”, where “array:0” is the dummy name of the first instance in the training or test data.
- In both `model_file` and `sys_output` files, **(`ci`, `pi`) pairs should be sorted by the spelling of `ci` alphabetically**. For instance, if the labels are `guns`, `mideast`, and `misc`, the line in `sys_output` should be `instanceName guns p1 mideast p2 misc p3`.
- `acc_file` shows the confusion matrix and the accuracy for the training and the test data (cf. **acc.ex**). In the confusion matrix, `a[i][j]` is the number of instances where the truth is class `i`, and the system output is class `j`.
- As always, `model.ex`, `sys.ex`, and `acc.ex` in the `examples/` directory are NOT gold standard. These files were created just to show you the format of the files.
- For your reference, `example_output/` shows the first 10 lines of the `sys_output` file if you build DT using `train.vectors.txt` and `test.vectors.txt` under `examples/` with the options `max_depth=4` and `min_gain=0.1` as specified in the `readme.txt` file.

Run **build\_dt.sh** with **train.vectors.txt** as the training data and **test.vectors.txt** as the test data:

---

<sup>1</sup>The depth of the root is 0, the depth of its children is 1, and so on.

- Fill out Table 2 (where min\_gain is set to 0) and Table 3 (where min\_gain is set to 0.1).
- submit model\_file, sys\_output, acc\_file produced by running  
`build_dt.sh train.vectors.txt test.vectors.txt 4 0.1 model_file sys_output > acc_file`

Table 2: Your decision tree results when min\_gain=0

| Depth | Training accuracy | Test accuracy | CPU time (in minutes) |
|-------|-------------------|---------------|-----------------------|
| 1     |                   |               |                       |
| 2     |                   |               |                       |
| 4     |                   |               |                       |
| 10    |                   |               |                       |
| 20    |                   |               |                       |
| 50    |                   |               |                       |

Table 3: Your decision tree results when min\_gain=0.1

| Depth | Training accuracy | Test accuracy | CPU time (in minutes) |
|-------|-------------------|---------------|-----------------------|
| 1     |                   |               |                       |
| 2     |                   |               |                       |
| 4     |                   |               |                       |
| 10    |                   |               |                       |
| 20    |                   |               |                       |
| 50    |                   |               |                       |

**Q4 (5 points):** Slide #12 of class2.DT.pdf shows a DT: f1 and f2 are two features; f1 is in [-20, 30]; f2 is in [-10, 30].  $L_i$  ( $i=1, \dots, 7$ ) represents a leaf node. Each leaf node corresponds to a rectangle in a 2-dimension space, where f1 is the x-axis and f2 is the y-axis. Draw a graph that shows the boundary of the seven rectangles in this 2-dimension space.

**Q5 (5 “free” points):** If you are not familiar with Patas or Condor submit, please go over the condor information at <https://wiki.ling.washington.edu/bin/view.cgi/Main/HowToUseCondor> and <https://wiki.ling.washington.edu/bin/view.cgi/Main/CondorClusterHomepage>. You can run condor submit for the code in Q3. We will use condor submit for many assignments later.

**Submission:** Submit the following to Canvas:

- Your note file *readme.(txt | pdf)* that includes your answers to Q1-Q4, and any notes that you want the TA to read.

- hw.tar.gz that includes all the files specified in `dropbox/21-22/572/hw2/submit-file-list`, plus any source code (and binary code) used by the shell scripts.
- Make sure that you run **check\_hw2.sh** before submitting your hw.tar.gz.
- No need to submit anything for Q5.